

# model training code

April 12, 2024

```
[1]: import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
from torchvision.datasets import CIFAR10
from torchsummary import summary
import matplotlib.pyplot as plt
```

```
[2]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")
```

Using device: cuda

```
[3]: transform_train = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ColorJitter(brightness=0.3, contrast=0.3, saturation=0.3, hue=0.
↪1),
    transforms.RandomRotation(degrees=20),
    transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
])

transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
])

trainset = CIFAR10(root='./data', train=True, download=True,
↪transform=transform_train)
trainloader = DataLoader(trainset, batch_size=128, shuffle=True, num_workers=2)

testset = CIFAR10(root='./data', train=False, download=True,
↪transform=transform_test)
```

```
testloader = DataLoader(testset, batch_size=100, shuffle=False, num_workers=2)
```

Files already downloaded and verified

Files already downloaded and verified

```
[4]: class BasicBlock(nn.Module):
    expansion = 1

    def __init__(self, in_planes, planes, stride=1):
        super(BasicBlock, self).__init__()
        self.conv1 = nn.Conv2d(in_planes, planes, kernel_size=3, stride=stride,
        ↪padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(planes)
        self.conv2 = nn.Conv2d(planes, planes, kernel_size=3, stride=1,
        ↪padding=1, bias=False)
        self.bn2 = nn.BatchNorm2d(planes)
        self.shortcut = nn.Sequential()
        if stride != 1 or in_planes != self.expansion*planes:
            self.shortcut = nn.Sequential(
                nn.Conv2d(in_planes, self.expansion*planes, kernel_size=1,
        ↪stride=stride, bias=False),
                nn.BatchNorm2d(self.expansion*planes)
            )

    def forward(self, x):
        out = torch.relu(self.bn1(self.conv1(x)))
        out = self.bn2(self.conv2(out))
        out += self.shortcut(x)
        out = torch.relu(out)
        return out
```

```
[5]: class ModifiedResNet(nn.Module):
    def __init__(self, block, num_blocks, num_classes=10):
        super(ModifiedResNet, self).__init__()
        self.in_planes = 64 # Increased number of initial channels

        self.conv1 = nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=1,
        ↪bias=False) # Increased initial channels
        self.bn1 = nn.BatchNorm2d(64)
        self.layer1 = self._make_layer(block, 64, num_blocks[0], stride=1)
        self.layer2 = self._make_layer(block, 128, num_blocks[1], stride=2) #
        ↪Increased channels
        self.layer3 = self._make_layer(block, 256, num_blocks[2], stride=2) #
        ↪Increased channels
        self.linear = nn.Linear(256*block.expansion, num_classes)

    def _make_layer(self, block, planes, num_blocks, stride):
```

```

        strides = [stride] + [1]*(num_blocks-1)
        layers = []
        for stride in strides:
            layers.append(block(self.in_planes, planes, stride))
            self.in_planes = planes * block.expansion
        return nn.Sequential(*layers)

    def forward(self, x):
        out = torch.relu(self.bn1(self.conv1(x)))
        out = self.layer1(out)
        out = self.layer2(out)
        out = self.layer3(out)
        out = nn.functional.avg_pool2d(out, 8)
        out = out.view(out.size(0), -1)
        out = self.linear(out)
        return out

```

```

[6]: def ResNetModel():
        return ModifiedResNet(ResidualBlock, [2,2,2])

    def count_parameters(model):
        return sum(p.numel() for p in model.parameters() if p.requires_grad)

```

```

[7]: def train_model(model, criterion, optimizer, scheduler, num_epochs=10):
        train_losses = []
        test_losses = []
        train_accs = []
        test_accs = []

        model.to(device)

        for epoch in range(num_epochs):
            model.train()
            running_loss = 0.0
            correct = 0
            total = 0
            for i, data in enumerate(trainloader, 0):
                inputs, labels = data[0].to(device), data[1].to(device)
                optimizer.zero_grad()

                outputs = model(inputs)
                loss = criterion(outputs, labels)
                loss.backward()
                optimizer.step()

                running_loss += loss.item()
                _, predicted = outputs.max(1)

```

```

        total += labels.size(0)
        correct += predicted.eq(labels).sum().item()

    train_loss = running_loss / len(trainloader)
    train_acc = correct / total
    train_losses.append(train_loss)
    train_accs.append(train_acc)

    model.eval()
    correct = 0
    total = 0
    test_loss = 0.0
    with torch.no_grad():
        for data in testloader:
            inputs, labels = data[0].to(device), data[1].to(device)
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            test_loss += loss.item()
            _, predicted = outputs.max(1)
            total += labels.size(0)
            correct += predicted.eq(labels).sum().item()

    test_loss /= len(testloader)
    test_acc = correct / total
    test_losses.append(test_loss)
    test_accs.append(test_acc)

    print(f'Epoch [{epoch+1}/{num_epochs}], Train Loss: {train_loss:.4f},
↪Train Acc: {train_acc:.4f}, Test Loss: {test_loss:.4f}, Test Acc: {test_acc:.
↪4f}')
    scheduler.step(test_loss)

plt.figure(figsize=(10, 5))
plt.plot(train_losses, label='Train Loss')
plt.plot(test_losses, label='Test Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Testing Loss Curves')
plt.legend()
plt.show()

plt.figure(figsize=(10, 5))
plt.plot(train_accs, label='Train Accuracy')
plt.plot(test_accs, label='Test Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training and Testing Accuracy Curves')

```

```
plt.legend()
plt.show()

print('Final Test Accuracy: {:.4f}'.format(test_acc))
```

```
[8]: model = ModifiedResNet(BasicBlock, [2, 2, 2, 2])
model.to(device)
```

```
[8]: ModifiedResNet(
  (conv1): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
    bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
    track_running_stats=True)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
        bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
        bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
      (shortcut): Sequential()
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
        bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
        bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
      (shortcut): Sequential()
    )
  )
  (layer2): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
        1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
        1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
```

```

        (shortcut): Sequential(
          (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (1): BasicBlock(
        (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (shortcut): Sequential()
      )
    )
    (layer3): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (shortcut): Sequential(
          (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (1): BasicBlock(
        (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (shortcut): Sequential()
      )
    )
    (linear): Linear(in_features=256, out_features=10, bias=True)

```

)

```
[9]: print(summary(model, (3, 32, 32)))
```

```
-----
      Layer (type)              Output Shape          Param #
=====
          Conv2d-1              [-1, 64, 32, 32]          1,728
        BatchNorm2d-2          [-1, 64, 32, 32]           128
          Conv2d-3              [-1, 64, 32, 32]        36,864
        BatchNorm2d-4          [-1, 64, 32, 32]           128
          Conv2d-5              [-1, 64, 32, 32]        36,864
        BatchNorm2d-6          [-1, 64, 32, 32]           128
        BasicBlock-7           [-1, 64, 32, 32]            0
          Conv2d-8              [-1, 64, 32, 32]        36,864
        BatchNorm2d-9          [-1, 64, 32, 32]           128
          Conv2d-10             [-1, 64, 32, 32]        36,864
        BatchNorm2d-11         [-1, 64, 32, 32]           128
        BasicBlock-12          [-1, 64, 32, 32]            0
          Conv2d-13            [-1, 128, 16, 16]        73,728
        BatchNorm2d-14         [-1, 128, 16, 16]          256
          Conv2d-15            [-1, 128, 16, 16]       147,456
        BatchNorm2d-16         [-1, 128, 16, 16]          256
          Conv2d-17            [-1, 128, 16, 16]         8,192
        BatchNorm2d-18         [-1, 128, 16, 16]          256
        BasicBlock-19          [-1, 128, 16, 16]            0
          Conv2d-20            [-1, 128, 16, 16]       147,456
        BatchNorm2d-21         [-1, 128, 16, 16]          256
          Conv2d-22            [-1, 128, 16, 16]       147,456
        BatchNorm2d-23         [-1, 128, 16, 16]          256
        BasicBlock-24          [-1, 128, 16, 16]            0
          Conv2d-25            [-1, 256, 8, 8]       294,912
        BatchNorm2d-26         [-1, 256, 8, 8]           512
          Conv2d-27            [-1, 256, 8, 8]       589,824
        BatchNorm2d-28         [-1, 256, 8, 8]           512
          Conv2d-29            [-1, 256, 8, 8]       32,768
        BatchNorm2d-30         [-1, 256, 8, 8]           512
        BasicBlock-31          [-1, 256, 8, 8]            0
          Conv2d-32            [-1, 256, 8, 8]       589,824
        BatchNorm2d-33         [-1, 256, 8, 8]           512
          Conv2d-34            [-1, 256, 8, 8]       589,824
        BatchNorm2d-35         [-1, 256, 8, 8]           512
        BasicBlock-36          [-1, 256, 8, 8]            0
          Linear-37             [-1, 10]                2,570
=====
Total params: 2,777,674
Trainable params: 2,777,674
```

Non-trainable params: 0

-----  
Input size (MB): 0.01  
Forward/backward pass size (MB): 10.50  
Params size (MB): 10.60  
Estimated Total Size (MB): 21.11  
-----

None

```
[10]: criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-4)
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min',
↪ factor=0.1, patience=5, verbose=True)

train_model(model, criterion, optimizer, scheduler, num_epochs=50)

torch.save(model.state_dict(), 'modified_resnet_cifar10_model_50_epochs.pth')
print('Model saved successfully!')
```

/home/nat/.local/lib/python3.10/site-packages/torch/optim/lr\_scheduler.py:28:  
UserWarning: The verbose parameter is deprecated. Please use get\_last\_lr() to  
access the learning rate.


warnings.warn("The verbose parameter is deprecated. Please use get\_last\_lr() "

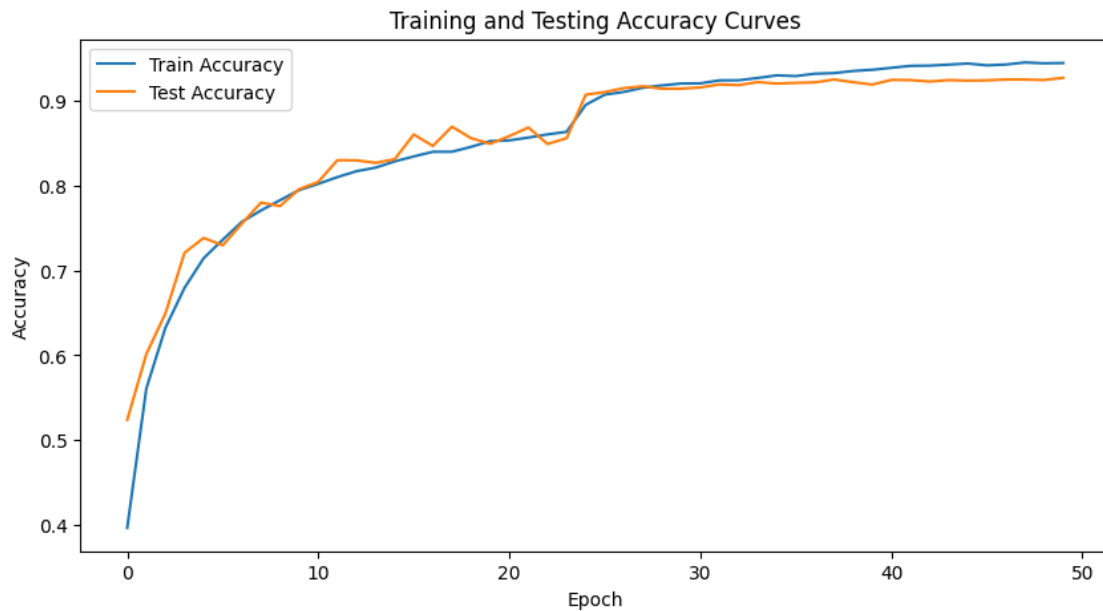
Epoch [1/50], Train Loss: 1.6236, Train Acc: 0.3969, Test Loss: 1.3303, Test  
Acc: 0.5241  
Epoch [2/50], Train Loss: 1.2241, Train Acc: 0.5609, Test Loss: 1.1386, Test  
Acc: 0.6015  
Epoch [3/50], Train Loss: 1.0363, Train Acc: 0.6325, Test Loss: 1.0775, Test  
Acc: 0.6495  
Epoch [4/50], Train Loss: 0.9151, Train Acc: 0.6794, Test Loss: 0.8460, Test  
Acc: 0.7207  
Epoch [5/50], Train Loss: 0.8173, Train Acc: 0.7143, Test Loss: 0.7502, Test  
Acc: 0.7382  
Epoch [6/50], Train Loss: 0.7521, Train Acc: 0.7361, Test Loss: 0.8804, Test  
Acc: 0.7295  
Epoch [7/50], Train Loss: 0.7022, Train Acc: 0.7571, Test Loss: 0.7500, Test  
Acc: 0.7550  
Epoch [8/50], Train Loss: 0.6616, Train Acc: 0.7705, Test Loss: 0.6930, Test  
Acc: 0.7798  
Epoch [9/50], Train Loss: 0.6251, Train Acc: 0.7826, Test Loss: 0.7053, Test  
Acc: 0.7757  
Epoch [10/50], Train Loss: 0.5946, Train Acc: 0.7945, Test Loss: 0.6230, Test  
Acc: 0.7956  
Epoch [11/50], Train Loss: 0.5725, Train Acc: 0.8018, Test Loss: 0.6164, Test  
Acc: 0.8041  
Epoch [12/50], Train Loss: 0.5459, Train Acc: 0.8098, Test Loss: 0.5196, Test  
Acc: 0.8297



Epoch [13/50], Train Loss: 0.5278, Train Acc: 0.8167, Test Loss: 0.5433, Test Acc: 0.8295  
Epoch [14/50], Train Loss: 0.5142, Train Acc: 0.8209, Test Loss: 0.5171, Test Acc: 0.8266  
Epoch [15/50], Train Loss: 0.4949, Train Acc: 0.8283, Test Loss: 0.5021, Test Acc: 0.8308  
Epoch [16/50], Train Loss: 0.4823, Train Acc: 0.8341, Test Loss: 0.4316, Test Acc: 0.8600  
Epoch [17/50], Train Loss: 0.4652, Train Acc: 0.8396, Test Loss: 0.4874, Test Acc: 0.8464  
Epoch [18/50], Train Loss: 0.4572, Train Acc: 0.8396, Test Loss: 0.3801, Test Acc: 0.8692  
Epoch [19/50], Train Loss: 0.4456, Train Acc: 0.8454, Test Loss: 0.4302, Test Acc: 0.8556  
Epoch [20/50], Train Loss: 0.4317, Train Acc: 0.8522, Test Loss: 0.4859, Test Acc: 0.8490  
Epoch [21/50], Train Loss: 0.4263, Train Acc: 0.8530, Test Loss: 0.4352, Test Acc: 0.8580  
Epoch [22/50], Train Loss: 0.4131, Train Acc: 0.8563, Test Loss: 0.4186, Test Acc: 0.8682  
Epoch [23/50], Train Loss: 0.4087, Train Acc: 0.8600, Test Loss: 0.4664, Test Acc: 0.8488  
Epoch [24/50], Train Loss: 0.4000, Train Acc: 0.8631, Test Loss: 0.4439, Test Acc: 0.8555  
Epoch [25/50], Train Loss: 0.3064, Train Acc: 0.8949, Test Loss: 0.2646, Test Acc: 0.9068  
Epoch [26/50], Train Loss: 0.2732, Train Acc: 0.9069, Test Loss: 0.2601, Test Acc: 0.9099  
Epoch [27/50], Train Loss: 0.2568, Train Acc: 0.9101, Test Loss: 0.2533, Test Acc: 0.9145  
Epoch [28/50], Train Loss: 0.2473, Train Acc: 0.9150, Test Loss: 0.2498, Test Acc: 0.9167  
Epoch [29/50], Train Loss: 0.2390, Train Acc: 0.9178, Test Loss: 0.2512, Test Acc: 0.9138  
Epoch [30/50], Train Loss: 0.2290, Train Acc: 0.9199, Test Loss: 0.2533, Test Acc: 0.9139  
Epoch [31/50], Train Loss: 0.2311, Train Acc: 0.9201, Test Loss: 0.2535, Test Acc: 0.9153  
Epoch [32/50], Train Loss: 0.2217, Train Acc: 0.9236, Test Loss: 0.2454, Test Acc: 0.9189  
Epoch [33/50], Train Loss: 0.2191, Train Acc: 0.9238, Test Loss: 0.2454, Test Acc: 0.9181  
Epoch [34/50], Train Loss: 0.2105, Train Acc: 0.9264, Test Loss: 0.2405, Test Acc: 0.9216  
Epoch [35/50], Train Loss: 0.2030, Train Acc: 0.9296, Test Loss: 0.2473, Test Acc: 0.9199  
Epoch [36/50], Train Loss: 0.2012, Train Acc: 0.9287, Test Loss: 0.2462, Test Acc: 0.9207

Epoch [37/50], Train Loss: 0.1970, Train Acc: 0.9315, Test Loss: 0.2410, Test Acc: 0.9213  
Epoch [38/50], Train Loss: 0.1929, Train Acc: 0.9322, Test Loss: 0.2428, Test Acc: 0.9247  
Epoch [39/50], Train Loss: 0.1882, Train Acc: 0.9347, Test Loss: 0.2414, Test Acc: 0.9215  
Epoch [40/50], Train Loss: 0.1827, Train Acc: 0.9361, Test Loss: 0.2532, Test Acc: 0.9187  
Epoch [41/50], Train Loss: 0.1760, Train Acc: 0.9385, Test Loss: 0.2372, Test Acc: 0.9242  
Epoch [42/50], Train Loss: 0.1705, Train Acc: 0.9407, Test Loss: 0.2371, Test Acc: 0.9240  
Epoch [43/50], Train Loss: 0.1681, Train Acc: 0.9410, Test Loss: 0.2355, Test Acc: 0.9223  
Epoch [44/50], Train Loss: 0.1648, Train Acc: 0.9422, Test Loss: 0.2356, Test Acc: 0.9240  
Epoch [45/50], Train Loss: 0.1629, Train Acc: 0.9435, Test Loss: 0.2349, Test Acc: 0.9234  
Epoch [46/50], Train Loss: 0.1667, Train Acc: 0.9414, Test Loss: 0.2366, Test Acc: 0.9237  
Epoch [47/50], Train Loss: 0.1637, Train Acc: 0.9423, Test Loss: 0.2337, Test Acc: 0.9247  
Epoch [48/50], Train Loss: 0.1605, Train Acc: 0.9447, Test Loss: 0.2347, Test Acc: 0.9246  
Epoch [49/50], Train Loss: 0.1638, Train Acc: 0.9437, Test Loss: 0.2338, Test Acc: 0.9241  
Epoch [50/50], Train Loss: 0.1611, Train Acc: 0.9441, Test Loss: 0.2318, Test Acc: 0.9267





Final Test Accuracy: 0.9267

Model saved successfully!

The model is trained using the Cross Entropy Loss function and optimized with the Adam optimizer, with a learning rate of 0.001 and weight decay set to 1e-4. Additionally, a learning rate scheduler, ReduceLROnPlateau, is employed to adjust the learning rate based on the validation loss.

During training, the model undergoes 50 epochs, during which the training and testing loss and accuracy are evaluated and displayed at each epoch. Both training and testing accuracies demonstrate a steady increase over the epochs, indicating effective learning. The loss and accuracy curves are plotted, depicting the gradual decrease in both training and testing loss alongside the increase in accuracy, affirming the model's learning progress.

```
[11]: def evaluate_model(model, dataloader):
    model.eval()
    correct = 0
    total = 0
    with torch.no_grad():
        for data in dataloader:
            inputs, labels = data[0].to(device), data[1].to(device)
            outputs = model(inputs)
            _, predicted = torch.max(outputs, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
    accuracy = correct / total
    print(f'Accuracy on the dataset: {100 * accuracy:.2f}%)')
```

```
evaluate_model(model, testloader)
```

Accuracy on the dataset: 92.67%